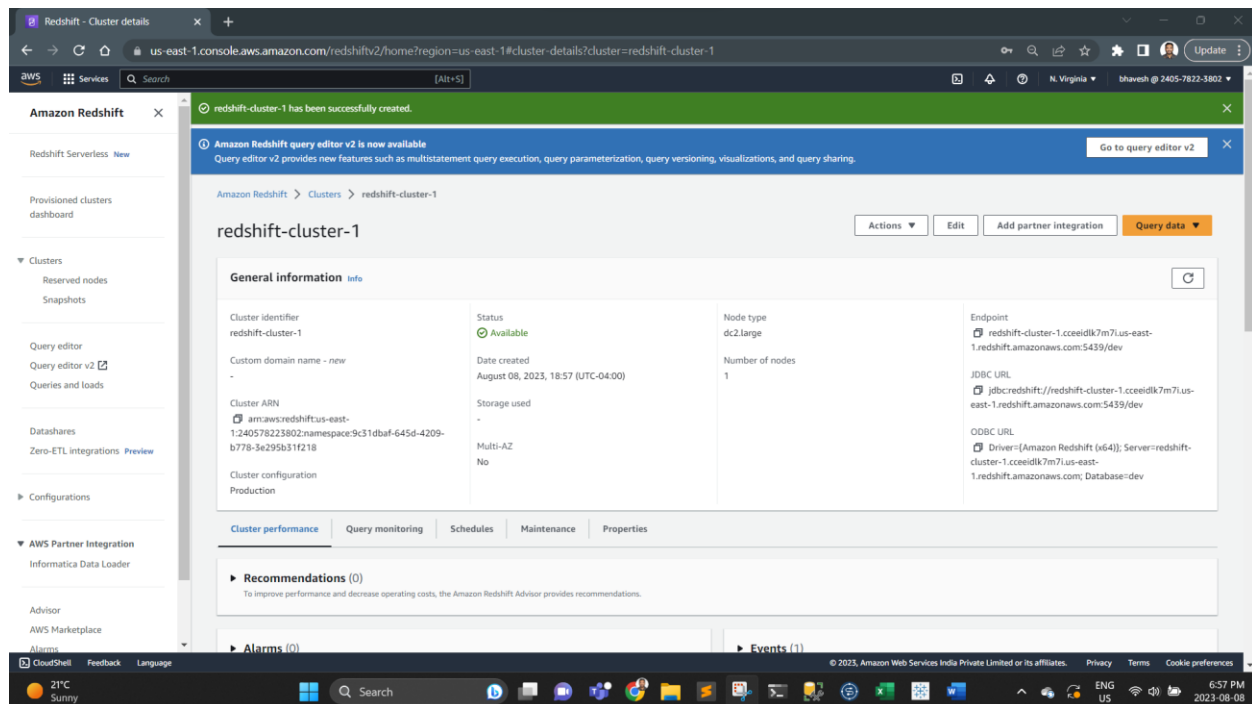


Prerequisites

Cluster and configure setup for Amazon Redshift ML administration.

Perform one-time cluster setup for Amazon Redshift ML

1. Creating a Redshift Cluster



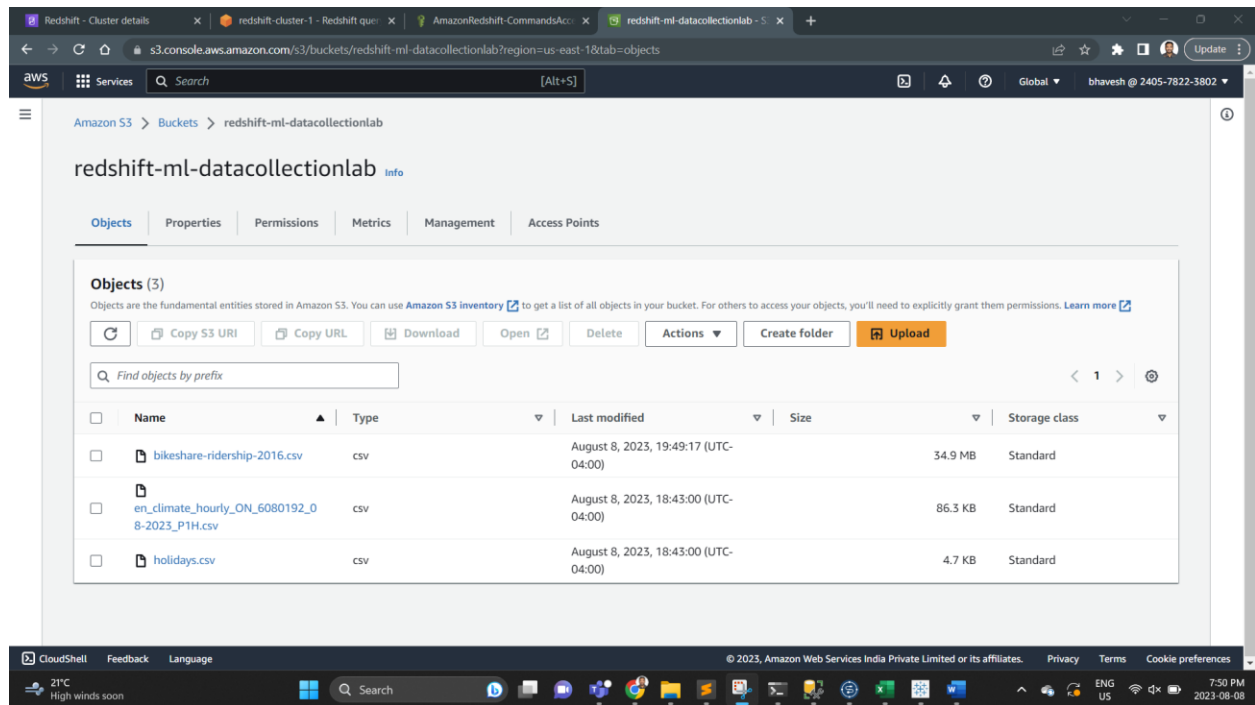
Created a Redshift Cluster and Assigned an IAM Role that has the **AmazonRedshiftAllCommandsFullAccess** policy automatically attached.

A simple operation allows us to create an IAM role with **AmazonS3FullAccess** and **AmazonSageMakerFullAccess** policies for use with Amazon Redshift ML.

The new IAM role that you create allows Amazon Redshift to copy, load, query, and analyze data from Amazon resources in your IAM account.

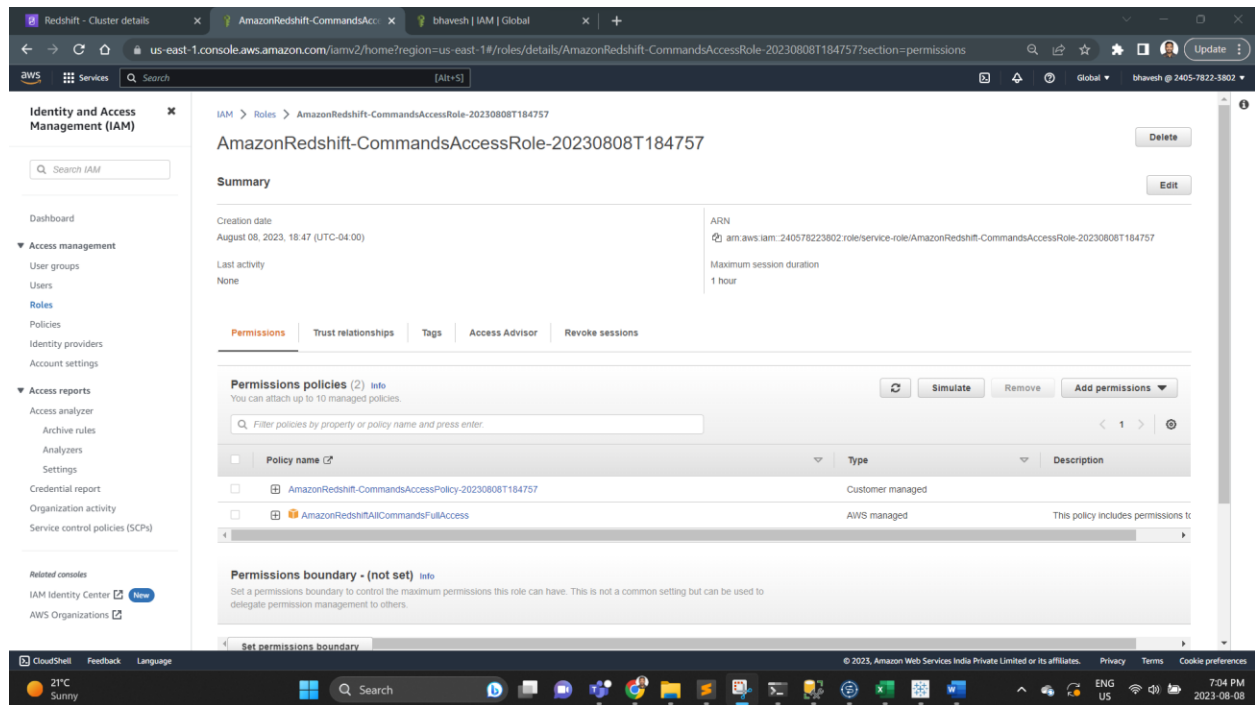
2. Create an S3 bucket and add all the data csv files to the bucket.

Data Collection and Curation – Machine Learning Lab – Final Project



- I choose a Specific Amazon S3 bucket : redshift-ml-datacollectionlab to specify that the IAM role being created has permission to access this bucket.


Data Collection and Curation – Machine Learning Lab – Final Project



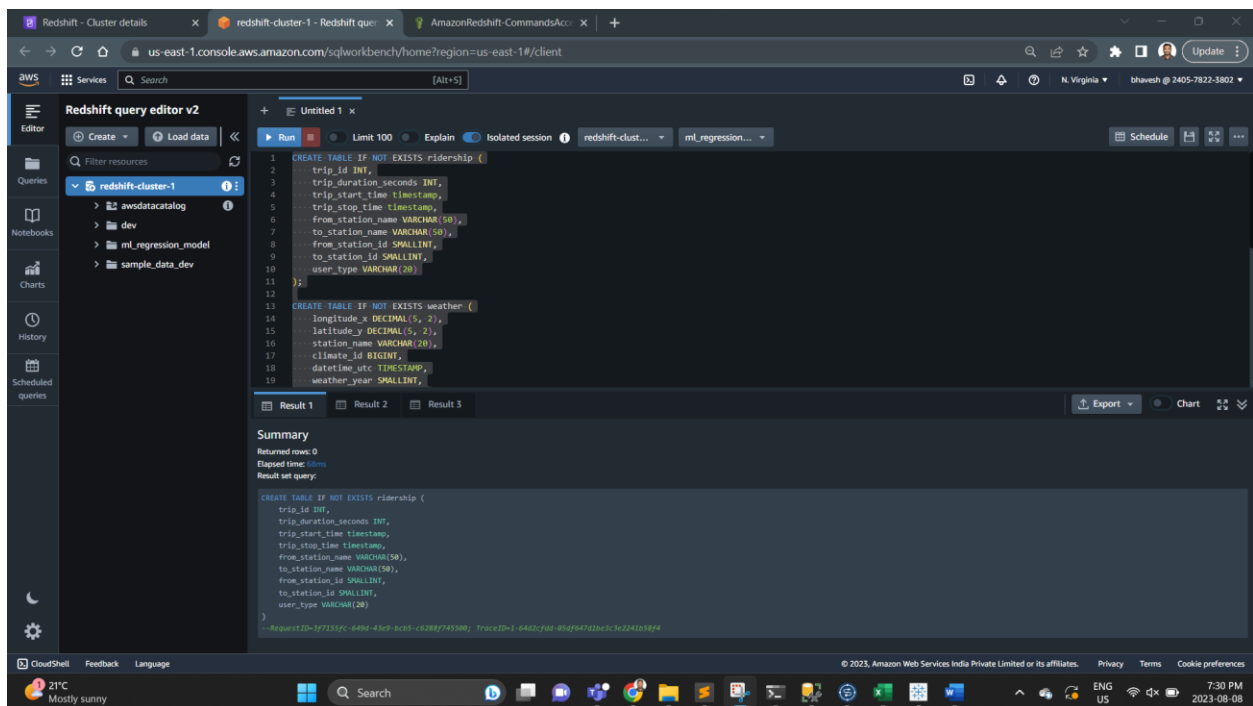
Step 1: Load the data from Amazon S3 to Amazon Redshift

Create a database inside the cluster to following creating tables there after.

Data Collection and Curation – Machine Learning Lab – Final Project



Select the specific database created and create the respective tables that are required to import data from.



Data Collection and Curation – Machine Learning Lab – Final Project

Load the records into the database tables created

The screenshot shows the AWS Redshift query editor interface. On the left, the 'Tables' section under 'public' shows the 'weather' table. The table schema is displayed below:

Field	Type	NULL	CMP
longitude_x	numeric(5,2)	NULL	az64
latitude_y	numeric(5,2)	NULL	az64
station_name	character varying(20)	NULL	lzo
climate_id	bigint	NULL	az64
datetime_utc	timestamp without time zone	NULL	az64
weather_year	smallint	NULL	az64
weather_month	smallint	NULL	az64
weather_day	smallint	NULL	az64
time_utc	character varying(5)	NULL	lzo

The main query editor shows a successful COPY query result:

```
1 COPY weather
2 FROM
3 's3://redshift-ml-datacollectionlab/en_climate_hourly_06-2023_PSH.csv'
4 IAM_ROLE default
5 FORMAT CSV
6 IGNOREHEADER 1
7 DATEFORMAT 'auto'
8 TIMEFORMAT 'auto'
9 REGION 'us-east-1';
```

The 'Summary' section indicates: 'Load into table 'weather' completed, 744 record(s) loaded successfully.' The 'Returned rows: 0' and 'Elapsed time: 18.7s' are also shown.

The screenshot shows the AWS Redshift query editor interface. On the left, the 'Tables' section under 'public' shows the 'holiday' table. The table schema is displayed below:

Field	Type	NULL	CMP
holiday_date	date	NULL	az64
description	character varying(100)	NULL	lzo

The main query editor shows a successful COPY query result:

```
13 FROM
14 's3://redshift-ml-datacollectionlab/bikeshare-ridership-2016.tsv'
15 IAM_ROLE default
16 IGNOREHEADER 1
17 DATEFORMAT 'auto'
18 TIMEFORMAT 'auto'
19 DELIMITER 't'
20 REGION 'us-east-1';
21
22 COPY holiday
23 FROM
24 's3://redshift-ml-datacollectionlab/holidays.csv'
25 IAM_ROLE default
26 FORMAT CSV
27 IGNOREHEADER 1
28 DATEFORMAT 'auto'
29 TIMEFORMAT 'auto'
30 REGION 'us-east-1';
31
```

The 'Summary' section indicates: 'Load into table 'holiday' completed, 198 record(s) loaded successfully.' The 'Returned rows: 0' and 'Elapsed time: 10.6s' are also shown.